




```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression , RANSACRegressor , TheilSenRegressor , Ridge , Lasso , ElasticNet , SGDRegressor
from sklearn.metrics import r2_score
from sklearn.svm import SVR
from sklearn.model_selection import train_test_split , cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from sklearn.ensemble import RandomForestRegressor , AdaBoostRegressor , GradientBoostingRegressor
from xgboost import XGBRegressor
from scipy.stats.mstats import winsorize
```

```
data = pd.read_csv('data.csv')
```

```
data.head()
```


	site area	structure type	water consumption	recycling rate	utilisation rate	air qality index	issue reolution time	resident count	electricity cost	 
0	1360	Mixed-use	2519.0	69	52	188	1	72	1420.0	
1	4272	Mixed-use	2324.0	50	76	165	65	261	3298.0	
2	3592	Mixed-use	2701.0	20	94	198	39	117	3115.0	
3	966	Residential	1000.0	13	60	74	3	35	1575.0	
4	4926	Residential	5990.0	23	65	32	57	185	4301.0	

Next steps:


 [View recommended plots](#)

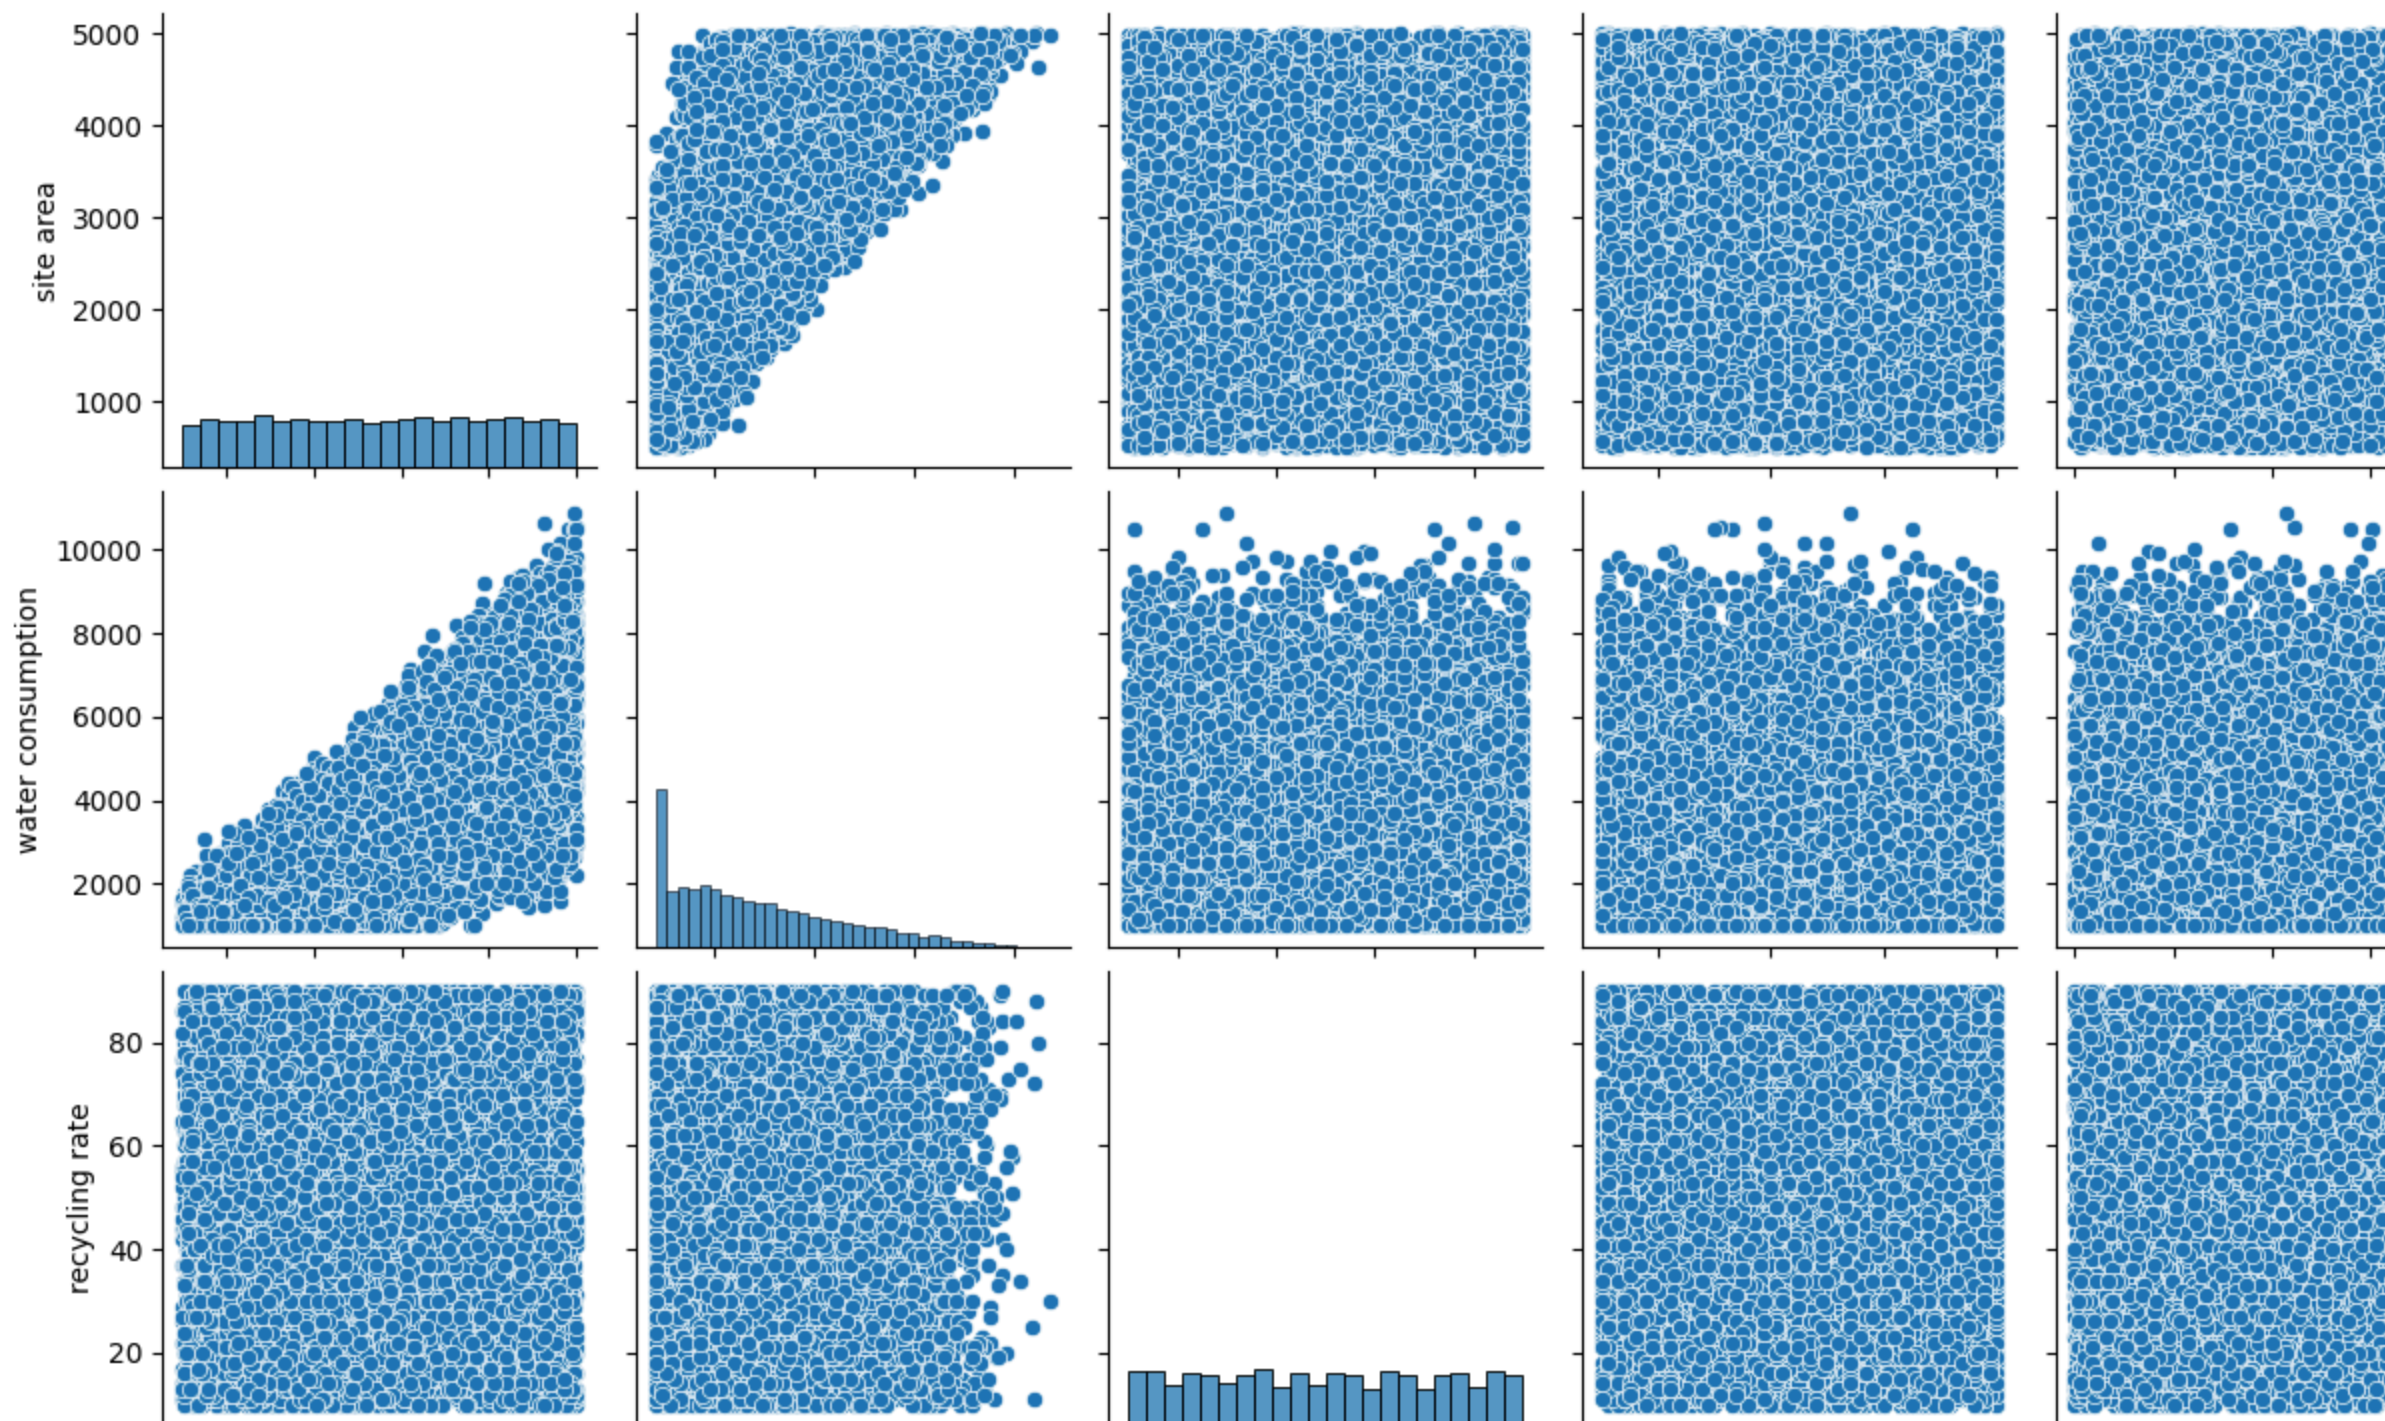
[New interactive sheet](#)

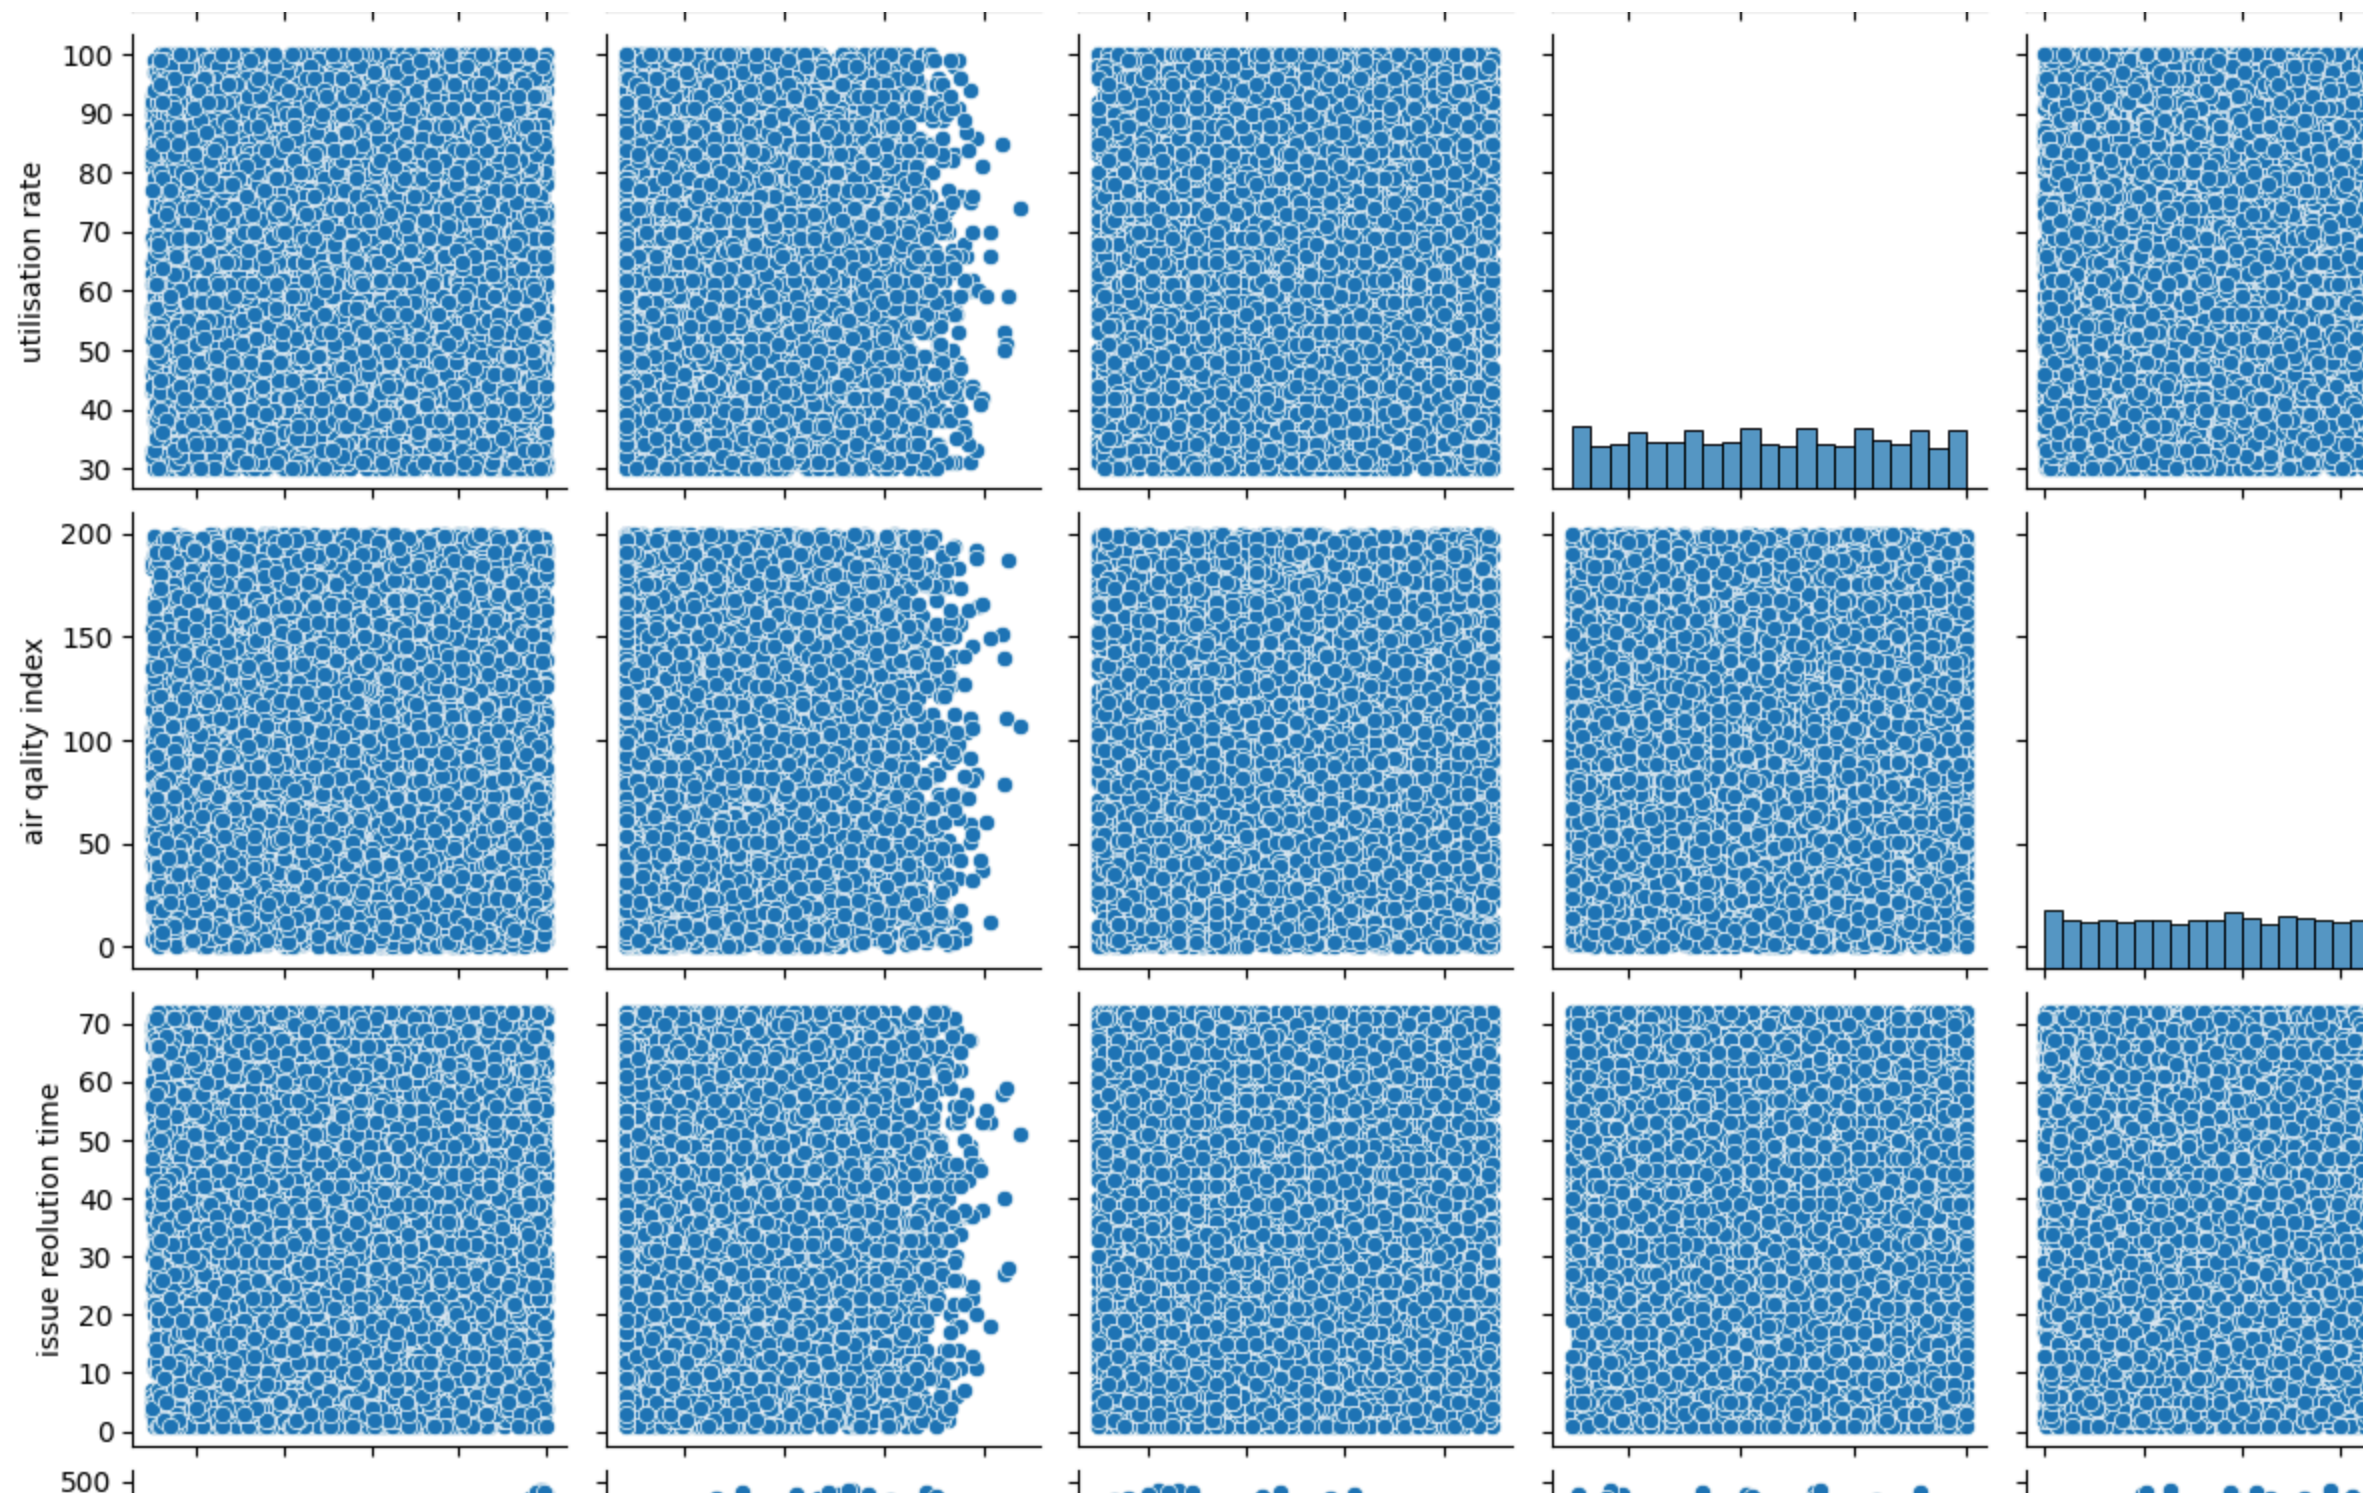
```
data.isnull().sum()
```


		0
<hr/>		
	site area	0
	structure type	0
	water consumption	0
	recycling rate	0
	utilisation rate	0
	air quality index	0
	issue reolution time	0
	resident count	0
	electricity cost	0
	dtype: int64	

```
sns.pairplot(data)
```

 <seaborn.axisgrid.PairGrid at 0x7bcd20757e90>







```
data.drop('structure type' , axis = 1 , inplace = True)
```

```
data.shape
```


```
➡ (10000, 8)
```

```
fig, ax = plt.subplots(4,2, figsize=(12,8))
axes_ = [axes_row for axes in ax for axes_row in axes]
```


```
for i, j in enumerate(data.columns):
    g = sns.boxplot(x = data[j], ax = axes_[i])
    g.set_title(j)
plt.tight_layout()
```




site area




water consumption




recycling rate




utilisation rate




air quality index




issue resolution time




resident count



electricity cost





```
col = data.columns
```

```
def check_outliers(data, col):
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5*IQR
    upper_bound = Q3 + 1.5*IQR

    outliers = data[col][(data[col] < lower_bound) | (data[col] > upper_bound)]


    print("Outliers Report")

    print(f"The total number of outliers in Data: {len(outliers)}")

    plt.figure(figsize=(10,10))
    plt.subplot(211)
    plt.plot(data[col])
    plt.title(col + " with Outliers")
    plt.scatter(x=outliers.index, y=outliers.values, marker="X", color='r', s=100)

    plt.subplot(212)
    plt.title(col + " after removing Extreme Values")
    filter_data = data[col][~(data[col].isin(outliers))]
    sns.boxplot(filter_data)
```

```
data.dtypes
```

	0
site area	int64
water consumption	float64
recycling rate	int64
utilisation rate	int64
air qality index	int64
issue reolution time	int64
resident count	int64
electricity cost	float64
dtype: object	


```
check_outliers(data, col[6])
```




Outliers Report


The total number of outliers in Data: 289

resident count with Outliers



resident count after removing Extreme Values






```
data_copy = data.copy()
```


```
def apply_winsorize(data, col):  
    winsorize(data[col], limits= [0.005, 0.005], inplace=True)  
for i, j in data.items():  
    apply_winsorize(data_copy, i)  
fig, ax = plt.subplots(4,2, figsize=(12,8))  
axes_ = [axes_row for axes in ax for axes_row in axes]  
  
for i, j in enumerate(data.columns):  
    g = sns.boxplot(x = data_copy[j], ax = axes_[i])  
    g.set_title(j)  
    plt.tight_layout()
```




site area




water consumption




recycling rate




utilisation rate




air quality index




issue resolution time



resident count



electricity cost





```
data_copy.dtypes
```

	0
site area	int64
water consumption	float64
recycling rate	int64
utilisation rate	int64
air qality index	int64
issue reolution time	int64
resident count	int64
electricity cost	float64


dtype: object

```
X = data_copy.iloc[: , data_copy.columns != 'electricity cost']
```

```
Y = data_copy['electricity cost']
```

```
sns.displot(Y , kind='kde')
```

 <seaborn.axisgrid.FacetGrid at 0x7bcd287bc810>



```
X_train , X_test ,y_train , y_test = train_test_split(X , Y , test_size=0.3 , random_state=0)
```

```
lr_pipeline = Pipeline([("scaler", StandardScaler()), ("linear_regression", LinearRegression())])
rb_pipeline = Pipeline([("scaler", StandardScaler()), ("robust_regression", RANSACRegressor(random_state=42))])
theil_pipeline = Pipeline([("scaler", StandardScaler()), ("theil_regressor", TheilSenRegressor(random_state=42))])
ridge_pipeline = Pipeline([("scaler", StandardScaler()), ("ridge_regressor", Ridge(random_state = 42))])
lasso_pipeline = Pipeline([("scaler", StandardScaler()), ("lasso_regressor", Lasso(random_state = 42))])
elastic_pipeline = Pipeline([("scaler", StandardScaler()), ("elastic_net", ElasticNet(random_state = 42))])
random_forest_pipeline = Pipeline([("scaler", StandardScaler()), ("randomforest_regression", RandomForestRegressor(random_state = 42))])
```

```

xgboost_pipeline = Pipeline([("scaler", StandardScaler()), ("xgboost_regression", XGBRegressor())])
adaboost_pipeline = Pipeline([("scaler", StandardScaler()), ("adaboost_regression", AdaBoostRegressor(random_state = 42))])
gradient_pipeline = Pipeline([("scaler", StandardScaler()), ("gradientboost_regression", GradientBoostingRegressor(random_state = 42))])
lightgbm_pipeline = Pipeline([("scaler", StandardScaler()), ("lightgbm_regression", LGBMRegressor(random_state = 42))])
decisiontree_pipeline = Pipeline([("scaler", StandardScaler()), ("decisiontree_regression", DecisionTreeRegressor(random_state = 42))])
knn_pipeline = Pipeline([("scaler", StandardScaler()), ("knn_regression", KNeighborsRegressor())])
sgc_pipeline = Pipeline([("scaler", StandardScaler()), ("sgd_regression", SGDRegressor(random_state = 42))])

```

```

pipelines = [lr_pipeline, rb_pipeline, theil_pipeline, ridge_pipeline, lasso_pipeline, elastic_pipeline,
             random_forest_pipeline, xgboost_pipeline, adaboost_pipeline, gradient_pipeline, lightgbm_pipeline,
             decisiontree_pipeline, knn_pipeline, sgc_pipeline]

```

```

pipe_dictionary = pipe_dict = {0: "Linear Regression", 1: "Robust", 2: "Theil Sen", 3: "Ridge", 4: "Lasso", 5: "ElasticNet",
                               6: "RandomForest", 7: "XGBoost", 8: "Adaboost", 9: "GradientBoost", 10: "LightGBM",
                               11: "Decision Tree", 12: "KNN", 13: "SGD"}

```

```

for i , pipe in enumerate(pipelines):
    score = cross_val_score(pipe , X , Y , cv=10)
    print(pipe_dictionary[i] , score.mean())

```



Show hidden output

```

import numpy as np
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
from lightgbm import LGBMRegressor

class ModelTuner:
    def __init__(self, model_name, search_type='grid', random_state=42):
        self.model_name = model_name.lower()
        self.search_type = search_type
        self.random_state = random_state
        self.model = None
        self.param_grid = None
        self.best_estimator = None

```

```

self.best_params = None
self.best_score = None

self._init_model_and_params()

def _init_model_and_params(self):
    if self.model_name == 'random_forest':
        self.model = RandomForestRegressor(random_state=self.random_state)
        self.param_grid = {
            'n_estimators': [100, 200],
            'max_depth': [None, 10, 20],
            'min_samples_split': [2, 5],
            'min_samples_leaf': [1, 2],
            'max_features': ['sqrt', 'log2']
        }
    elif self.model_name == 'lightgbm':
        self.model = LGBMRegressor(random_state=self.random_state)
        self.param_grid = {
            'n_estimators': [100, 200],
            'learning_rate': [0.01, 0.05, 0.1],
            'max_depth': [-1, 5, 10],
            'num_leaves': [31, 50],
            'min_child_samples': [10, 20],
            'subsample': [0.7, 0.9],
            'colsample_bytree': [0.7, 0.9]
        }
    else:
        raise ValueError("Unsupported model. Use 'random_forest' or 'lightgbm'.")

def fit(self, X, y, test_size=0.2, cv=5, n_jobs=-1):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=self.random_state)
    sc = StandardScaler()
    X_train = sc.fit_transform(X_train)
    X_test = sc.transform(X_test)

    if self.search_type == 'grid':
        search = GridSearchCV(self.model, self.param_grid, cv=cv, scoring='r2', n_jobs=n_jobs)
    else:
        from sklearn.model_selection import RandomizedSearchCV

```



```
search = RandomizedSearchCV(self.model, self.param_grid, cv=cv, scoring='r2', n_jobs=n_jobs, n_iter=30, random_state=self.random_state)
```